


Assignment 2



[COSC2767] Systems Deployment and Operations

Team Pall

Do Tung Lam
S3963286

Nguyen Quoc An
S3938278

Truong Quang Bao
Loc
S3965528

Hoang Thai Phuc
S3978081

Submission Date: January 17th, 2025

Lecturer: Dr. Tom Huynh

Contents

I. CI/CD Pipeline Solution	2
1. Problem statement	2
2. Proposed Solution.....	2
3. Tools and technologies	2
II. CI/CD Pipeline Solution	2
4. Problem statement	2
5. Proposed Solution.....	3
6. Tools and technologies	3
7. Pipeline features	3
III. Pipeline Component	4
IV. Pipeline Advanced Component	18
8. Infrastructure as Code (IaC)	18
9. Monitoring and Alerting	19
a. Prometheus	19
b. Grafana	21
10. Canary Deployment	22
V. Application Flow (Diagram)	24
VI. Known Bugs/ Problems.....	25
VII. Pipeline Self-Evaluation	25
VIII. Project Responsibility	26
IX. Conclusion.....	26
X. Reference	27
XI. Appendix.....	27

I. CI/CD Pipeline Solution

1. Problem statement

Dr. Tom, a web developer managing the RMIT E-commerce store, struggles to efficiently deploy updates due to his limited DevOps experience and the absence of a CI/CD pipeline.

2. Proposed Solution

To address Tom's challenges, we propose CI/CD as the core solution to automate deployments, reduce errors, and enhance reliability. Thanks to continuous integration, the code will be integrated smoothly, while continuous delivery will provide safe, staged releases with rollback capabilities. Stability would be improved because of automated testing and monitoring, gaining real-time system visibility. Structured workflow improves collaboration, scalability, and efficiency, hence making deployments seamless and predictable.

3. Tools and technologies

- **GitHub:** A platform that helps us to manage and control the version of the RMIT store [1]. This platform will be the starting point of the CI/CD, when there is any commit to change the source code, it will trigger the pipeline.
- **Jenkins:** This is an open-source automation server that allows us to implement all types of automation tasks from building, testing to deploying our products [2]. We will use this platform for complex workflows like automation testing.
- **GitHub Action:** This is a CI tool offered by GitHub and it is only used within GitHub [3]. This tool will automatically execute our workflow and also simplify some processes for Jenkins like trigger on pull request with file pattern.
- **AWS CloudFormation:** This is a service that helps us model and set up AWS resources and provides more time for applications operating in AWS rather than dealing with resource management [4]. It is used to bring up the test environment that simulates the production environment then we can shut down and eliminate that environment after testing for cost efficiency.
- **Docker:** Docker is a software platform that enables one to build, test, and deploy applications quickly. Docker packages software into standardized units called containers; these contain everything the software needs to run, such as libraries, system tools, code, and even runtime [5]. With Docker, we can ensure that RMIT store can run stably across different environments and simplify the deployment process when we just need to install docker on different environments.
- **MongoDB Atlas:** MongoDB Atlas is a fully managed cloud database that will handle all the complexities of deploying, managing, and healing your deployments on the cloud service provider of your choice: AWS, Azure, and GCP [6]. This is our main choice to deploy our database.
- **Nginx:** This is an HTTP web server, reverse proxy, content cache, load balancer, TCP/UDP proxy server, and mail proxy server [7]. In our CI/CD, Nginx acts as a reverse proxy, it provides zero-downtime deployment and ensures secured access of applications.
- **Elastic Load Balancer:** Elastic Load Balancing (ELB) automatically distributes incoming application traffic across multiple backend servers or front-end servers, Load Balancer enables the scalability, availability and reliability of the website [8].
- **Kubernetes:** Kubernetes is an open-source, portable, extensible platform for managing containerized workloads and services which facilitates both declarative configuration and automation [9]. Kubernetes allows our CI/CD to deploy automated scaling, rollbacks, load balancing, service discovery, configuration management, monitoring, and zero-downtime updates and we can also apply canary deployment with K8s.
- **Prometheus:** An open-source monitoring system with a dimensional data model, flexible query language, efficient time series database and modern alerting approach [10]. In our CI/CD pipeline this tool will help us to monitor the website.
- **Grafana:** This is the open-source analytics & monitoring solution for every database [11]. In our CI/CD, we use Grafana to integrate with Prometheus for visualizing time-series metrics that allow engineers to detect anomalies, optimize performance, and troubleshoot issues efficiently.

II. CI/CD Pipeline Solution

4. Problem statement

Dr. Tom, a web developer managing the RMIT E-commerce store, struggles to efficiently deploy updates due to his limited DevOps experience and the absence of a CI/CD pipeline.

5. Proposed Solution

To address Tom's challenges, we propose CI/CD as the core solution to automate deployments, reduce errors, and enhance reliability. Thanks to continuous integration, the code will be integrated smoothly, while continuous delivery will provide safe, staged releases with rollback capabilities. Stability would be improved because of automated testing and monitoring, gaining real-time system visibility. Structured workflow improves collaboration, scalability, and efficiency, hence making deployments seamless and predictable.

6. Tools and technologies

- **GitHub:** A platform that helps us to manage and control the version of the RMIT store [1]. This platform will be the starting point of the CI/CD, when there is any commit to change the source code, it will trigger the pipeline.
- **Jenkins:** This is an open-source automation server that allows us to implement all types of automation tasks from building, testing to deploying our products [2]. We will use this platform for complex workflows like automation testing.
- **GitHub Action:** This is a CI tool offered by GitHub and it is only used within GitHub [3]. This tool will automatically execute our workflow and also simplify some processes for Jenkins like trigger on pull request with file pattern.
- **AWS CloudFormation:** This is a service that helps us model and set up AWS resources and provides more time for applications operating in AWS rather than dealing with resource management [4]. It is used to bring up the test environment that simulates the production environment then we can shut down and eliminate that environment after testing for cost efficiency.
- **Docker:** Docker is a software platform that enables one to build, test, and deploy applications quickly. Docker packages software into standardized units called containers; these contain everything the software needs to run, such as libraries, system tools, code, and even runtime [5]. With Docker, we can ensure that RMIT store can run stably across different environments and simplify the deployment process when we just need to install docker on different environments.
- **MongoDB Atlas:** MongoDB Atlas is a fully managed cloud database that will handle all the complexities of deploying, managing, and healing your deployments on the cloud service provider of your choice: AWS, Azure, and GCP [6]. This is our main choice to deploy our database.
- **Nginx:** This is an HTTP web server, reverse proxy, content cache, load balancer, TCP/UDP proxy server, and mail proxy server [7]. In our CI/CD, Nginx acts as a reverse proxy, it provides zero-downtime deployment and ensures secured access of applications.
- **Elastic Load Balancer:** Elastic Load Balancing (ELB) automatically distributes incoming application traffic across multiple backend servers or front-end servers, Load Balancer enables the scalability, availability and reliability of the website [8].
- **Kubernetes:** Kubernetes is an open-source, portable, extensible platform for managing containerized workloads and services which facilitates both declarative configuration and automation [9]. Kubernetes allows our CI/CD to deploy automated scaling, rollbacks, load balancing, service discovery, configuration management, monitoring, and zero-downtime updates and we can also apply canary deployment with K8s.
- **Prometheus:** An open-source monitoring system with a dimensional data model, flexible query language, efficient time series database and modern alerting approach [10]. In our CI/CD pipeline this tool will help us to monitor the website.
- **Grafana:** This is the open-source analytics & monitoring solution for every database [11]. In our CI/CD, we use Grafana to integrate with Prometheus for visualizing time-series metrics that allow engineers to detect anomalies, optimize performance, and troubleshoot issues efficiently.

7. Pipeline features

- **Automated Code Integration:** Whenever a new pull request is created in GitHub, a build and test process is automatically triggered. It promotes continuous merging and validation, thereby simplifying change management and catching errors early in the development lifecycle. By integrating code in smaller batches, teams maintain a synchronized repository and reduce the likelihood of complex merge conflicts.
- **Automated Code Build and Testing:** Jenkins manages complex workflows that encompass unit tests, integration tests, and other quality checks. By automating these processes, Jenkins removes much of the human error that manual testing can introduce. This

consistent approach to validating every code change maintains the project’s stability and frees developers to focus on building new features rather than troubleshooting deployment issues.

- **Automated Test Environment Management:** Using AWS CloudFormation, the pipeline dynamically provisions and tears down AWS resources for testing, effectively mirroring the production environment. This ensures consistency in infrastructure setup and eliminates the risk of configuration drift between environments. Moreover, it is cost-efficient by automatically spinning down test environments when they are no longer needed.
- **Consistent Environment and Streamlined Deployments:** Docker packages the application and all its dependencies into lightweight containers, allowing the software to run uniformly across different environments. By encapsulating everything the application needs, Docker makes it easier to ensure predictable performance and reliability in development, testing, and production stages.
- **Traffic Scalability:** Elastic Load Balancing automatically distributes traffic across multiple server instances or containers, preventing any single resource from becoming a bottleneck. As the site’s traffic grows or changes, it scales resources accordingly to maintain optimal performance. This redundancy ensures the store remains available and responsive under varying loads.
- **Automated Deployment Orchestration and Management:** Kubernetes streamlines container deployments self-healing capabilities restart failed containers or replace them on different hosts if necessary. By supporting canary deployment strategy, Kubernetes reduces the risk of introducing errors into production and facilitates agile release cycles. Thanks to container orchestration and robust version control, the pipeline can quickly revert to a known stable version if a release introduces critical issues. This rollback process greatly reduces downtime and user disruption. A clear record of versions and their respective configurations ensures the team can isolate and correct problems with minimal impact on the live environment.
- **Real-time Monitoring and Alerting:** Prometheus collects real-time metrics and stores them in a time-series database, while Grafana provides dynamic dashboards that visualize these metrics. This integration enables proactive monitoring, giving the team immediate insight into any anomalies or performance overload. In addition. Grafana alerting rules which send alert notifications allow operators to respond quickly to potential issues before they escalate into outages.

III. Pipeline Component

For each completed main requirement, please provide description, configurations and screenshots as proofs.

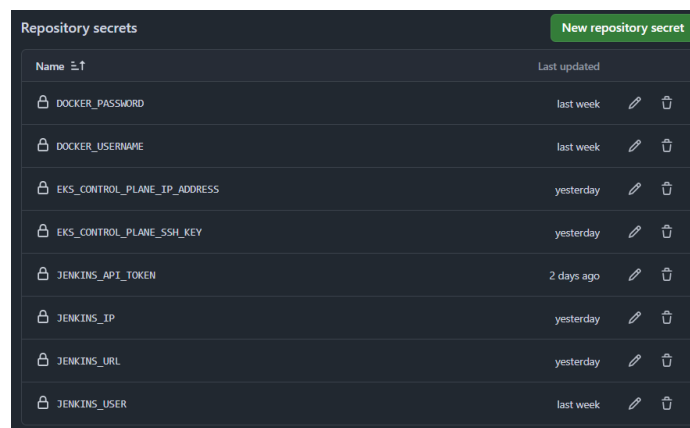
Explanation of the chosen settings and configurations if needed.

Screenshots are important to show the pipeline can pass required use cases.

Secret and Credentials Configuration

To allow the CI/CD pipeline to work securely, sensitive credentials must be properly configured and use within the workflow.

GitHub Repository Secrets



















Repository secrets		New repository secret
Name	Last updated	
DOCKER_PASSWORD	last week	 
DOCKER_USERNAME	last week	 
EKS_CONTROL_PLANE_IP_ADDRESS	yesterday	 
EKS_CONTROL_PLANE_SSH_KEY	yesterday	 
JENKINS_API_TOKEN	2 days ago	 
JENKINS_IP	yesterday	 
JENKINS_URL	yesterday	 
JENKINS_USER	last week	 

Figure 1 Repository Secrets

DOCKER_USERNAME & DOCKER_PASSWORD: These 2 two secrets store the credentials to login Docker registry.

JENKINS_IP: store the IP of Jenkins server help securely trigger the Jenkins job.

JENKINS_API_TOKEN: This secret will store the API token for credentials when call the API remotely.

JENKINS_URL: This secret stores the URL that the pipeline can trigger the Jenkins job.

JENKINS_USER: Store the username for authenticating Jenkins.

Credentials









T	P	Store i	Domain	ID	Name
		System	(global)	ssh-github-key	git (ssh-github-key)
		System	(global)	locoioioi-dockerhub	locoioioi/***** (locoioioi-dockerhub)
		System	(global)	ec2-ssh-vockey	ec2-user (ec2-ssh-vockey)
		System	(global)	eks-cluster-ip	EKS Cluster IP for Deployment

Figure 2 Jenkins credentials

ssh-github-key: A ssh key for authenticating with GitHub repositories.

locoioioi-dockerhub: Credentials for Docker Hub associated with the account `locoioioi`.

ec2-ssh-vockey: SSH private key to connect with Ec2 instance.

eks-cluster-ip: Stores the IP address of the EKS control plan server for deployment purposes.

General GitHub Action Workflow

a) Workflow for linting check with GitHub Action

```
1  name: Lint Check
2
3  on:
4    pull_request:
5      branches:
6        - main
7
8  jobs:
9    lint:
10     runs-on: ubuntu-latest
11
12     steps:
13       # Step 1: Check out the repository
14       - name: Check out code
15         uses: actions/checkout@v3
16
17       # Step 2: Set up Node.js
18       - name: Set up Node.js
19         uses: actions/setup-node@v3
20         with:
21           node-version: '17' # Adjust base
22
23       # Step 3: Install dependencies
24       - name: Install dependencies
25         run: npm install
26
27       # Step 4: Run linting
28       - name: Run linting
29         run: npm run lint
```

Eslint

Purpose: Perform eslint check for NodeJS and ReactJS code.

Name: Lint Check

Trigger On: This job is triggered on a pull request event to the main branch.

Jobs: This job will be run on the ubuntu-latest image.

Steps:

- Check out the code using the actions/checkout action.
- Set up Node.js version 17 using the actions/setup-node action.
- Install dependencies with the npm install command.
- Run linting using the npm run lint command.

Figure 3. GitHub Lint Check Job]

```

1  name: reviewdog
2  on: [pull_request]
3  jobs:
4    eslint:
5      name: runner / eslint
6      runs-on: ubuntu-latest
7      permissions:
8        contents: read
9        pull-requests: write
10     steps:
11       - name: Checkout code
12         uses: actions/checkout@v3
13
14       - name: Set up Node.js
15         uses: actions/setup-node@v3
16         with:
17           node-version: '17' # Update this to match your project
18
19       - name: Install dependencies
20         run: npm install
21
22       - name: Run ESLint
23         run: npx eslint . -f json -o eslint-report.json || true
24
25       - name: Run reviewdog
26         uses: reviewdog/action-eslint@v1
27         with:
28           reporter: github-pr-review
29           level: error
30           eslint_input: eslint-report.json
31           fail_on_error: false
32           filter_mode: nofilter

```

Figure 5 ReviewDog Lint

```

45  client-tests:
46    name: Client Unit Tests
47    runs-on: ubuntu-latest
48
49    steps:
50      # Step 1: Check out the code from the repository
51      - name: Checkout repository
52        uses: actions/checkout@v3
53
54      # Step 2: Set up Node.js
55      - name: Setup Node.js
56        uses: actions/setup-node@v3
57        with:
58          node-version: 20
59
60      # Step 3: Install client dependencies
61      - name: Install client dependencies
62        run: |
63          cd client
64          npm install
65          npm install --save-dev jest supertest
66
67      # Step 4: Run client tests if "test" script exists
68      - name: Run client tests
69        run: |
70          cd client
71          if npm run | grep -q "test"; then
72            echo "Running client tests..."
73            npm run test
74          else
75            echo "No test script found. Skipping client tests."
76          fi
77      env:
78        NODE_ENV: test

```

Figure 4 ReviewDog Lint

ReviewDog Lint

Name: Reviewdog

Purpose: Display the eslint result with ReviewDog Bot in comment and GitHub checks tab.

Trigger On: This job is triggered on a pull request event.

Jobs: This job will be run on the ubuntu-latest image.

- Check out the code using the actions/checkout action.
- Set up Node.js version 17 using the actions/setup-node action.
- Install dependencies with the npm install command.
- Run ESLint to generate a JSON report (eslint-report.json) using the npx eslint command.
- Use the reviewdog/action-eslint action to annotate ESLint errors on the pull request with the github-pr-review reporter, error level set to error, and nofilter mode enabled.

```

1  name: Unit and Integration Test CI
2  |
3  on:
4    pull_request:
5      branches:
6        - main
7    workflow_dispatch:
8
9  jobs:
10   server-tests:
11     name: Server Unit Tests
12     runs-on: ubuntu-latest
13
14     steps:
15       # Step 1: Check out the code from the repository
16       - name: Checkout repository
17         uses: actions/checkout@v3
18
19       # Step 2: Set up Node.js
20       - name: Setup Node.js
21         uses: actions/setup-node@v3
22         with:
23           node-version: 20
24
25       # Step 3: Install server dependencies
26       - name: Install server dependencies
27         run: |
28           cd server
29           npm install
30           npm install --save-dev jest supertest
31
32       # Step 4: Run server tests if "test" script exists
33       - name: Run server tests
34         run: |
35           cd server
36           if npm run | grep -q "test"; then
37             echo "Running server tests..."
38             npm run test
39           else
40             echo "No test script found. Skipping server tests."
41           fi
42       env:
43         NODE_ENV: test

```

Figure 7 Unit and Integration Test CI

```

1  name: Unit and Integration Test CI
2  |
3  on:
4    pull_request:
5      branches:
6        - main
7    workflow_dispatch:
8
9  jobs:
10   server-tests:
11     name: Server Unit Tests
12     runs-on: ubuntu-latest
13
14     steps:
15       # Step 1: Check out the code from the repository
16       - name: Checkout repository
17         uses: actions/checkout@v3
18
19       # Step 2: Set up Node.js
20       - name: Setup Node.js
21         uses: actions/setup-node@v3
22         with:
23           node-version: 20
24
25       # Step 3: Install server dependencies
26       - name: Install server dependencies
27         run: |
28           cd server
29           npm install
30           npm install --save-dev jest supertest
31
32       # Step 4: Run server tests if "test" script exists
33       - name: Run server tests
34         run: |
35           cd server
36           if npm run | grep -q "test"; then
37             echo "Running server tests..."
38             npm run test
39           else
40             echo "No test script found. Skipping server tests."
41           fi
42       env:
43         NODE_ENV: test

```

Figure 6 Unit and Integration Test CI

b) Workflow for Unit Test and Integration Test

Name: Unit and Integration Test CI

Trigger On: This job is triggered on a pull request event to the main branch and can also be triggered manually using workflow dispatch.

1. Server Unit Tests:

- Runs on the ubuntu-latest image.
- Steps:
 - Check out the repository using the actions/checkout action.
 - Set up Node.js version 20 using the actions/setup-node action.
 - Install server dependencies, including Jest and Supertest, using npm install commands in the server directory.
 - Run server tests if a test script exists in the server directory. The NODE_ENV is set

to test.

2. Client Unit Tests:

- Runs on the ubuntu-latest image.
- Steps:
 - Check out the repository using the actions/checkout action.
 - Set up Node.js version 20 using the actions/setup-node action.
 - Install client dependencies, including Jest and Supertest, using npm install commands in the client directory.
 - Run client tests if a test script exists in the client directory. The NODE_ENV is set to test.

Continuous Integration

Frontend pipeline

```
jobs:
  build-push-trigger:
    runs-on: ubuntu-latest

    steps:
      # Step 1: Checkout the code from the repository
      - name: Checkout code
        uses: actions/checkout@v3

      # Step 2: Log in to Docker Hub
      - name: Log in to Docker Hub
        uses: docker/login-action@v2
        with:
          username: ${ secrets.DOCKER_USERNAME }
          password: ${ secrets.DOCKER_PASSWORD }

      # Step 3: Build and Push Docker E2E Testing Image
      - name: Build Docker E2E Testing image
        run: |
          docker build -t ${ secrets.DOCKER_USERNAME }/e2e:latest -f ./client/Dockerfile.e2e ./client

      - name: Push Docker E2E Testing image
        run: |
          docker push ${ secrets.DOCKER_USERNAME }/e2e:latest

      # Step 4: Trigger Jenkins Job (using Secret for Jenkins IP)
      - name: Trigger Jenkins Job
        run: |
          curl -X POST "http://${ secrets.JENKINS_IP }:8080/job/fe-ci/buildWithParameters?branch=${ github.head_ref }" \
            --user ${ secrets.JENKINS_USER }:${ secrets.JENKINS_API_TOKEN } \
            --header "Content-Type: application/json"

      - name: Install jq
        run: sudo apt-get install jq

      # Step 5: Wait for Jenkins Job to Complete (using Secret for Jenkins IP)
      - name: Wait for Jenkins Job
        id: wait_jenkins
        run: |
          JENKINS_URL="http://${ secrets.JENKINS_IP }:8080"
          JOB_NAME="fe-ci"
          USER="${ secrets.JENKINS_USER }"
          API_TOKEN="${ secrets.JENKINS_API_TOKEN }"
          BUILD_URL="${ JENKINS_URL / job / $JOB_NAME / lastBuild / api / json }"
          STATUS="IN_PROGRESS"

          while [ "$STATUS" != "SUCCESS" ] && [ "$STATUS" != "FAILURE" ]; do
            echo "Checking Jenkins job status..."
            sleep 30
            STATUS=$(curl -s -u $USER:$API_TOKEN $BUILD_URL | jq -r '.result')
            echo "Current status: $STATUS"
          done

          if [ "$STATUS" == "SUCCESS" ]; then
            echo "Jenkins job completed successfully!"
          else
            echo "Jenkins job failed. Exiting with error."
            exit 1
          fi
```

Figure 8 GitHub Action - Frontend CI for Pull Request

b) Developing a workflow for Pull Request on Frontend with GitHub Action

Name: Jenkins CI PR FE

Trigger On: This job is trigger on pull request event to main when the changes are pushed to folder client.

Jobs: This job will be run on the ubuntu-latest image.

Steps:

1. Login to docker hub with the DOCKERHUB_USERNAME and DOCKERHUB_PASSWORD secrets.
2. Build the e2e test image and then push it to docker hub with the tag "latest".
3. Trigger Jenkins Job named "fe-ci" with the repository secrets JENKINS IP to ensure the security of the workflow.
4. Wait for Jenkins job to finish and update the GitHub action job base on that status.

c) Configure Jenkins job "fe-ci" that is trigger by the Jenkins PR CI

- Ssh to the Jenkins server on AWS then login as root user.
- Run the command "service jenkins start" to start Jenkins.
- Access the Jenkins via port 8080 with the public ipv4 of the ec2 instance.
- Create a new Pipeline item with the name of "fe-ci"



Figure 10 Set up SCM for the Jenkins job

Provide the ssh URL of the repository on the Repository URL field.

Select appropriate credentials to allow access to private repository.

Specify the correct path to the Jenkinsfile in the Script Path field.

```
pipeline {
    agent any

    environment {
        STACK_NAME = "test-env-aws-deployment-stack"
        INSTANCE_IP = ""
    }

    parameters {
        string(defaultValue: 'main', description: 'Branch to build', name: 'branch')
    }

    stages {
        stage('Checkout Code') {
            steps {
                script {
                    checkout([
                        $class: 'GitSCM',
                        branches: [[name: 'refs/heads/${params.branch}']],
                        userRemoteConfigs: [[
                            credentialsId: 'ssh-github-key',
                            url: 'git@github.com:RMIT-Vietnam-Teaching/cosc2767-assignment-2-group-2024c-pall.git'
                        ]]
                    ])
                }
            }
        }
    }
}
```

Figure 11 Jenkinsfile - checkout code

- This is the Jenkinsfile that the pipeline will execute, first it will get a parameter call branch like what we configured on the pipeline.
- The first step of this pipeline is to checkout to the branch from the parameter with the appropriate credentials.

```
stage('Trigger CloudFormation') {
    steps {
        script {
            sh """
            aws cloudformation deploy \
            --stack-name ${STACK_NAME} \
            --template-file cloudFormation/create-test-server-template.yaml \
            --capabilities CAPABILITY_IAM
            aws cloudformation wait stack-create-complete --stack-name ${STACK_NAME}
            """
        }
    }
}

stage('Retrieve Instance IP') {
    steps {
        script {
            INSTANCE_IP = sh(
                script: """
                aws cloudformation describe-stacks \
                --stack-name ${STACK_NAME} \
                --query "Stacks[0].Outputs[?OutputKey=='PublicIP'].OutputValue" \
                --output text
                """,
                returnStdout: true
            ).trim()
            echo "EC2 Instance IP: ${INSTANCE_IP}"
        }
    }
}
```

Figure 13 Jenkinsfile - Run AWS CloudFormation template

- In the next stage, the pipeline will run the AWS CloudFormation template to create a new ec2 instance for testing purposes. CloudFormation template will be described detailed in the advanced section below.
- After successfully launching the new ec2 instance, it will retrieve the INSTANCE_IP and store it within the environment variable.

Figure 14 Jenkinsfile - config the testing environment

Figure 15 Run AWS CloudFormation template

```
stage('Install Docker & Dependencies') {
    steps {
        withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariables: ['SSH_KEY'])]) {
            sh """
            ssh -i ${SSH_KEY} -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} \
            "sudo yum install -y docker git nmap-ncat && \
            "nmap"
            sudo systemctl start docker && \
            sudo systemctl enable docker && \
            sudo usermod -s /usr/sbin/dmccs ec2-user && \
            "usermod"
            docker --version && \
            sudo curl -L https://github.com/docker/compose/releases/download/1.22.0/docker-compose-$(uname -s)-$(uname -m) -o /usr/local/bin/docker-compose && \
            sudo chmod +x /usr/local/bin/docker-compose && \
            docker-compose version"
            """
        }
    }
}

stage('Copy Source Code to EC2') {
    steps {
        withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariables: ['SSH_KEY'])]) {
            retry(3) {
                sh """
                scp -i ${SSH_KEY} -o StrictHostKeyChecking=no -r . ec2-user@${INSTANCE_IP}:/home/ec2-user/project
                """
            }
        }
    }
}
```

Figure 16 Install dependency

- In the test environment, first the pipeline will **install all necessary tools** to run testing like docker, and docker-compose
- After that, in the next stage, the pipeline will forward all the source code to the test environment.

```

stage('Deploy Application Using Docker Compose') {
  steps {
    withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
      sh """
      ssh -i \${SSH_KEY} -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} \
      =
      cd /home/ec2-user/project && \
      docker-compose up -d --build && \
      echo 'Waiting for containers to become healthy...' && \
      for i in $(seq 1 10); do \
      docker-compose ps | grep -q 'healthy' && break || sleep 5; \
      done || (echo 'Containers failed to become healthy' && exit 1)
      """
    }
  }
}

stage('Run E2E Tests') {
  steps {
    withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
      sh """
      ssh -i \${SSH_KEY} -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} \
      =
      sudo docker system prune -af && \
      docker run --name e2e-container -e TEST_URL=http://mlt-store-client-service \
      --network project_default localhost/e2e && \
      docker cp e2e-container:/e2e/results /home/ec2-user/cypress-reports
      """
    }
  }
}

stage('Retrieve Test Results') {
  steps {
    withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
      sh """
      scp -i \${SSH_KEY} -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP}:/home/ec2-user/cypress-reports/test-results*.xml ./cypress-reports/
      """
    }
  }
}

stage('Analyze Test Results') {
  steps {
    junit 'cypress-reports/test-results*.xml'
  }
}

```

Figure 17 Jenkinsfile - run test

- Figure 8 indicates that the Jenkinsfile will run the website with the `docker-compose up -d --build` command and then wait for it to be healthy.
- Next, the pipeline will **run end-to-end (E2E) tests** by executing a **Docker container** that runs Cypress tests against the deployed website.
- The test results will be copied from the container to the EC2 instance using `docker cp`.
- Afterward, the pipeline will **retrieve the test results** by securely copying them (`scp`) from the EC2 instance to the Jenkins workspace.
- Finally, the pipeline will **analyze the test results** using `junit`, allowing Jenkins to display test reports and determine the success or failure of the job.

```

post {
  success {
    echo 'Testing passed!'
  }
  failure {
    echo 'Testing failed!'
  }
  always {
    echo 'Cleaning up resources...'
    script {
      try {
        sh """
        aws cloudformation delete-stack --stack-name ${STACK_NAME}
        """
        echo 'CloudFormation stack deletion triggered.'
      } catch (Exception e) {
        echo "Failed to delete CloudFormation stack: ${e.message}"
      }
    }
    echo 'Job finished!'
  }
}

```

Figure 18 Jenkinsfile - post action

- After the pipeline is finished, it will always **clean up resources** like delete the **testing environment**.
- The pipeline will determine whether the job is a **success** or **failure** based on the stage **“Analyze test result”**.

```

describe("Product display in the shop test", () => {
  before(() => {
    cy.visit("/shop");
  });

  it("should display a list of products", () => {
    cy.get(".product-list").should("exist");

    cy.get(".product-container")
      .should("exist")
      .and("have.length.greaterThan", 0);
  });
});

describe("Product detail page test", () => {
  before(() => {
    cy.visit("/shop");
  });

  it("should navigate to the product detail page", () => {
    cy.get(".product-container").first().click();

    cy.url().should("include", "/product/");
  });

  after(() => {
    cy.go("back");
  });
});

```

```

describe("Filtering test", () => {
  before(() => {
    cy.visit("/shop");
  });

  it("Should display products high to low price", () => {
    cy.contains("Newest First").click();

    cy.contains("Price High to Low").click();

    cy.get(".product-list").should("exist");

    cy.wait(2000);
    cy.get(".product-container .price").then((priceElements) => {
      const prices = [...priceElements].map((el) => {
        parseFloat(el.innerText.replace("$", ""))
      });

      const sortedPrices = [...prices].sort((a, b) => b - a);
      expect(prices).to.deep.equal(sortedPrices);
    });
  });
});

describe("Login Page Tests", () => {
  beforeEach(() => {

```

Figure 19 E2e test

- This is the e2e test that the pipeline above will build image and run this test, and you can access this folder through client/cypress/e2e.
- It includes product display in the shop test, product detail page display test, filtering test, login test, add to cart functionality test, place order functionality test.

Backend pipeline

d) Workflow for Pull Request on Back-end with GitHub Action

Name: Jenkins CI PR BE

Trigger On: This job is triggered on a pull request event to the main branch when changes are pushed to the server folder.

Build and Push Docker Image:

- Check out the code from the pull request branch using the actions/checkout action.
- Log in to Docker Hub using the docker/login-action action with credentials from GitHub secrets.
- Extract the short commit SHA for tagging the Docker image.
- Build the Docker image for the server using the extracted commit SHA as the tag.
- Push the tagged Docker image to Docker Hub.

Trigger Jenkins Job:

- Define variables including Jenkins job name, URL, user, API token, and the Docker image tag.
- Obtain a Jenkins Crumb for secure API access.
- Trigger the Jenkins job with the Docker image tag as a parameter using an HTTP POST request.
- Check the HTTP response to confirm successful job triggering.

Wait for Jenkins Job Completion:

- Continuously check the Jenkins job status at regular intervals until it completes.
- Fetch and display Jenkins console logs upon completion.
- Exit with an error if the Jenkins job fails.

```

name: Jenkins CI PR BE

on:
  pull_request:
    branches:
      - main
    paths:
      - "server/**"

jobs:
  build-and-push:

```

```

40  run: |
41      # Define variables
42      JOB_NAME="be-ci"
43      JENKINS_URL="${{ secrets.JENKINS_URL }}"
44      USER="${{ secrets.JENKINS_USER }}"
45      API_TOKEN="${{ secrets.JENKINS_API_TOKEN }}"
46      echo "Image Tag: ${{ secrets.DOCKER_USERNAME }}/be-mern-server:${{ env.sha_short }}"
47      echo "PR Number: ${{ github.event.pull_request.number }}"
48
49      # Get Jenkins Crumb
50      CRUMB=$(curl -s --user "${USER}:${API_TOKEN}" "${JENKINS_URL}/crumbIssuer/api/json" | jq -r ".crumb")
51
52      # Trigger Jenkins job and get queue URL
53      # curl -s -X POST "${JENKINS_URL}/job/${JOB_NAME}/buildWithParameters" \
54      #   --user "${USER}:${API_TOKEN}" \
55      #   -H "Jenkins-Crumb:${CRUMB}" \
56      #   --data-urlencode "image_tag=${{ secrets.DOCKER_USERNAME }}/be-mern-server:${{ env.sha_short }}"
57
58      # Trigger Jenkins job and get HTTP response code
59      HTTP_RESPONSE=$(curl -o /dev/null -s -w "%{httpcode}" -X POST "${JENKINS_URL}/job/${JOB_NAME}/buildWithParameters" \
60        --user "${USER}:${API_TOKEN}" \
61        -H "Jenkins-Crumb:${CRUMB}" \
62        -d "image_tag=${{ secrets.DOCKER_USERNAME }}/be-mern-server:${{ env.sha_short }}")
63
64      # Check if the response is in the 400-405 range
65      if [[ "$HTTP_RESPONSE" =~ 400 || "$HTTP_RESPONSE" =~ 405 ]]; then
66        echo "Failed to trigger Jenkins job, HTTP response code: $HTTP_RESPONSE"
67        exit 1
68      fi
69
70      echo "Successfully triggered Jenkins job, HTTP response code: $HTTP_RESPONSE"
71
72      # Wait for Jenkins Job to Complete and Show Logs
73      - name: Wait for Jenkins Job
74      id: wait-jenkins
75      run: |
76        JOB_NAME="be-ci"
77        JENKINS_URL="${{ secrets.JENKINS_URL }}"
78        USER="${{ secrets.JENKINS_USER }}"
79        API_TOKEN="${{ secrets.JENKINS_API_TOKEN }}"
80        BUILD_URL="${JENKINS_URL}/job/${JOB_NAME}/lastBuild/api/json"
81        CONSOLE_OUTPUT_URL="${JENKINS_URL}/job/${JOB_NAME}/lastBuild/consoleText"
82        STATUS="IN PROGRESS"
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100

```

Figure 20 Jenkins CI PR BE Github Action Job

e) Trigger the Jenkins CI jobs with GitHub Action

- Repeat the steps to create a Jenkins Pipeline above, but change the string parameter to “image_tag”, and Script Path to “server/jenkins/Jenkinsfile”. Use the same credentials as the previous job for frontend CI.

☒ This project is parameterized ?

String Parameter ?

Name ?

Default Value ?

Description ?

Script Path ?

Figure 21 Jenkins CI BE Config

- This is the Jenkinsfile that the pipeline will execute, the GitHub Action jenkins-ci-be.yaml will trigger this file with input parameters equal to the extracted commit SHA.

```

1  pipeline {
2      agent any
3
4      environment {
5          STACK_NAME = "test-env-be-mern-deployment-stack"
6          INSTANCE_IP = ""
7      }
8
9      parameters {
10         string(name: "image_tag", defaultValue: "locoioio/be-mern-server:latest", description: "Docker image tag to deploy")
11     }
12
13     stages {
14         stage('Ensure Stack is Deleted Before Start') {
15             steps {
16                 script {
17                     echo "Checking for existing stack and ensuring deletion..."
18                     def stackStatus = sh(
19                         script: """
20                         aws cloudformation describe-stacks --stack-name "${STACK_NAME}" --query "Stacks[0].StackStatus" --output text || echo "STACK_DOES_NOT_EXIST"
21                         """,
22                         returnStdout: true
23                     ).trim()
24
25                     if (stackStatus != "STACK_DOES_NOT_EXIST") {
26                         echo "Stack exists with status: ${stackStatus}. Deleting stack..."
27                         sh """
28                         aws cloudformation delete-stack --stack-name "${STACK_NAME}"
29                         echo "Waiting for stack to delete..."
30                         aws cloudformation wait stack-delete-complete --stack-name "${STACK_NAME}"
31                         """,
32                         returnStdout: true
33                     )
34                     echo "No existing stack found. Continue with the pipeline..."
35                 }
36             }
37         }
38
39         stage('Trigger CloudFormation') {
40             steps {
41                 script {
42                     echo "Creating AWS Stack..."
43                     sh """
44                     aws cloudformation deploy \
45                         --stack-name "${STACK_NAME}" \
46                         --template-file cloudformation/create-test-server-be-template.yaml \
47                         --capabilities CAPABILITY_IAM
48                     """
49                 }
50             }
51         }
52
53         stage('Wait for Instance') {
54             steps {
55                 script {
56                     echo "Waiting for 1 minute to ensure the server is ready..."
57                     sh "sleep 60"
58                 }
59             }
60         }
61
62         stage('Retrieve Instance IP') {
63             steps {
64                 script {
65                     echo "Retrieving EC2 Instance IP..."
66                     INSTANCE_IP = sh(
67                         script: """
68                         aws cloudformation describe-stacks \
69                             --stack-name "${STACK_NAME}" \
70                             --query "Stacks[0].Outputs[?OutputKey=='PublicIP'].OutputValue" \
71                             --output text
72                         """,
73                         returnStdout: true
74                     ).trim()
75                     echo "EC2 Instance IP: ${INSTANCE_IP}"
76                 }
77             }
78         }
79     }
80 }

```

Figure 22 Jenkinsfile for CI BE

- This Jenkins job will first check if a cloudformation stack already exist or not according to the defined cloudformation template in the github repo, if yes, it will removed the stack before going on.
- If not exist, it will create the stack, which include an EC2 server, and other required resources, then retrieve the IP address.

```

181 stage('Install Dependencies') {
182     steps {
183         withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
184             sh """
185             ssh -i $SSH_KEY -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} "
186             echo 'Starting: Update system packages'
187             sudo yum update -y
188             echo 'Complete: Update system packages'
189
190             echo 'Starting: Install Docker'
191             sudo yum install -y docker
192             sudo service docker start
193             echo 'Complete: Install Docker'
194
195             echo 'Starting: Install Node.js using npm'
196             # Install npm
197             curl -o https://raw.githubusercontent.com/nvm-sh/nvm/v0.39.4/install.sh | bash
198
199             # Load nvm in this session
200             export NVM_DIR="/home/ec2-user/.nvm"
201             [ -s "/home/ec2-user/.nvm.sh" ] && \. "/home/ec2-user/.nvm.sh"
202             [ -s "/home/ec2-user/.bash_completion" ] && \. "/home/ec2-user/.bash_completion"
203
204             # Install and use the latest LTS version of Node.js
205             nvm install --lts
206             nvm use --lts
207             echo 'Complete: Install Node.js using npm'
208
209             echo 'Starting: Verify Node.js installation'
210             node -v
211             echo 'Complete: Verify Node.js installation'
212
213             echo 'Starting: Verify npm installation'
214             npm -v
215             echo 'Complete: Verify npm installation'
216
217             echo 'Install Dependencies Successfully'
218             =
219             """
220         }
221     }
222 }

```

```

126 stage('Install MongoDB') {
127     steps {
128         withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
129             sh """
130             ssh -i $SSH_KEY -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} "
131             # Add the MongoDB repository
132             echo '[mongodb-org-6.0]' | sudo tee /etc/yum.repos.d/mongodb-org-6.0.repo
133             echo 'name=MongoDB Repository' | sudo tee -a /etc/yum.repos.d/mongodb-org-6.0.repo
134             echo 'baseurl=https://repo.mongodb.org/yum/amazon/2023/mongodb-org-6.0/x86_64/' | sudo tee -a /etc/yum.repos.d/mongodb-org-6.0.repo
135             echo 'gpgcheck=1' | sudo tee -a /etc/yum.repos.d/mongodb-org-6.0.repo
136             echo 'enabled=1' | sudo tee -a /etc/yum.repos.d/mongodb-org-6.0.repo
137             echo 'gpgkey=https://www.mongodb.org/static/pgp/server-6.0.asc' | sudo tee -a /etc/yum.repos.d/mongodb-org-6.0.repo
138
139             # Install and Start MongoDB
140             sudo yum install -y mongodb-org
141             sudo systemctl start mongod
142             sudo systemctl enable mongod
143             mongod --version
144             sudo yum remove -y mongodb-mongosh
145             sudo yum install -y mongodb-mongosh-shared-openssl3
146             sudo yum install -y mongodb-mongosh
147             echo 'Install MongoDB Successfully'
148             =
149             """
150         }
151     }
152 }

```

Figure 23 Jenkinsfile for CI BE

- It and install them.
- will then SSH into the EC2 and install Docker, NodeJS, MongoDB

```

151 stage('Deploy and Seed MongoDB') {
152     steps {
153         withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
154             sh """
155             ssh -i $SSH_KEY -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} "
156             echo 'Pulling the docker image: $image_tag'
157             sudo docker pull $image_tag
158             echo 'Image pulled successfully'
159
160             # Stop and remove any existing container named be-mern-server
161             sudo docker stop be-mern-server || true
162             sudo docker rm be-mern-server || true
163
164             # Run a temporary container to seed the database
165             sudo docker run --rm \
166             --name mongo-test \
167             --network mongo-test \
168             --env MONGO_URI=mongodb://localhost:2702/mern-commerce \
169             $image_tag \
170             npm run seed:db admin@wtf.edu.cn mypassword
171
172             # Run the application container
173             sudo docker run -d --name be-mern-server -p 3000:3000 \
174             --name mongo-test \
175             --network mongo-test \
176             --env MONGO_URI=mongodb://localhost:2702/mern-commerce \
177             $image_tag
178
179             echo 'Deploy and Seed MongoDB Successfully'
180             =
181             """
182         }
183     }
184 }

```

```

188 stage('Health Check') {
189     steps {
190         withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
191             sh """
192             ssh -i $SSH_KEY -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} "
193             echo 'Waiting for the server to be ready...'
194             sleep 10
195             echo 'Testing healthcheck endpoint'
196             curl -v -X GET http://localhost:3000/healthcheck || (echo 'Healthcheck failed' && exit 1)
197             =
198             """
199         }
200     }
201 }
202
203 stage('MongoDB Health Check') {
204     steps {
205         withCredentials([sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')]) {
206             sh """
207             ssh -i $SSH_KEY -o StrictHostKeyChecking=no ec2-user@${INSTANCE_IP} "
208             mongosh --eval 'db.runCommand({ ping: 1 })' || (echo 'MongoDB ping failed' && exit 1)
209             =
210             """
211         }
212     }
213 }
214 }

```

```

216 post {
217     always {
218         script {
219             echo 'Cleaning up. Deleting AWS stack...'
220             sh """
221             aws cloudformation delete-stack --stack-name ${STACK_NAME}
222             aws cloudformation wait stack-delete-complete --stack-name ${STACK_NAME}
223             """
224         }
225         echo 'Pipeline completed.'
226     }
227     failure {
228         echo 'Pipeline failed!'
229     }
230 }
231 }

```

Figure 24 Jenkinsfile for CI BE

- It pull the docker image with the tag of the commit SHA, and run it. It also start the MongoDB database and seed it.
- Then it performs a healthcheck to port 3000 of the server container, and test the connection to the local MongoDB.
- If everything pass, or even if not it will cleanup the job by deleting all resources in the CloudFormation stack.

Continuous Deployment

Frontend pipeline

a) Set up Jenkins job to handle push event to main.

- Connect to Jenkins server and create new pipeline item with the name fe-deploy.
- Config the Jenkins job to trigger though webhooks of repository and read the correct jenkinsfile path.

Figure 25 Jenkins job - configuration

```
stage('Checkout Code') {
  steps {
    git branch: 'main', credentialsId: 'ssh-github-key', url: 'git@github.com:RMIT-Vietnam-Teaching/cpsc2767-assignment-2-group-2024c-pall.git'
  }
}

stage('Install Dependencies') {
  steps {
    dir('client') {
      sh 'npm ci'
    }
  }
}

stage('Build React App') {
  steps {
    dir('client') {
      sh 'npm run build'
    }
  }
}
```

- This is the first step of the deployment job for Frontend, after checkout to the main branch the pipeline will install dependency and build it.

Figure 26 Jenkinsfile - build react project

```
stage('Manage Docker Image') {
  steps {
    dir('client') {
      script {
        // Login to Docker registry
        docker.withRegistry('https://index.docker.io/v1/', 'localhost-dockerhub') {
          // Pull the existing latest image
          sh "docker pull ${IMAGE_NAME}:${IMAGE_VERSION}"

          // Tag the pulled image as 'stable'
          sh "docker tag ${IMAGE_NAME}:${IMAGE_VERSION} ${IMAGE_NAME}:${STABLE_TAG}"

          // Push the 'stable' image to Docker Hub
          sh "docker push ${IMAGE_NAME}:${STABLE_TAG}"

          // Build a new image tagged as 'latest'
          def customImage = docker.build("${IMAGE_NAME}:${IMAGE_VERSION}")
          // Push the new 'latest' image to Docker Hub
          customImage.push()
        }
      }
    }
  }
}

stage('Deploy to Kubernetes') {
  steps {
    script {
      withCredentials([
        string(credentialsId: 'eks-cluster-ip', variable: 'EKS_IP'),
        sshUserPrivateKey(credentialsId: 'ec2-ssh-vockey', keyFileVariable: 'SSH_KEY')
      ]) {
        sh """
        ssh -i ${SSH_KEY} -o StrictHostKeyChecking=no ec2-user@${EKS_IP} \
        "sudo -i sh -c 'whoami && kubectl rollout restart deployment rmit-store-client-canary'"
        """
      }
    }
  }
}
```

- After being built successfully, the pipeline will get the current latest tag from the registry and change it to stable tag, then it will push a new latest version of the frontend.
- In the last stage, the pipeline will ssh to the EKS control plan then exec kubectl rollout restart deployment rmit-store-client-canary to deploy.

Figure 27 Jenkinsfile - upadte image and rollout new deployment




Figure 28 CD results

- Figure 15 indicates the pipeline is completed successfully and rollout a new version for RMIT store.

Backend pipeline

b) Workflow to handle push event to main and PR closure

Name: Re-Tag and Push Latest Image to Docker Hub When PR is Merged

Trigger On: This job is triggered on pull request closures to the main branch that modify the server folder, only if the pull request is merged.

Update Latest Tag:

- Check out the merged pull request code using the actions/checkout action
- Log in to Docker Hub with credentials from GitHub secrets

- Extract the short commit SHA for tagging Docker images
- Pull the current latest image, tag it as stable, and push the stable tag to Docker Hub
- Pull the new image by commit SHA, tag it as latest, and push the latest tag to Docker Hub
- Set up an SSH key using the EKS control plane SSH key from GitHub secrets
- SSH into the EKS control plane and restart the deployment for the rmit-store-server-canary Kubernetes deployment

```

1 name: Re-Tag and Push Latest Image to Docker Hub when PR is Merged
2
3 on:
4   pull_request:
5     branches:
6       - main
7   paths:
8     - 'server/**'
9   types: [closed]
10
11 jobs:
12   update-latest-tag:
13     if: ${{ github.event.pull_request.merged }}
14     runs-on: ubuntu-latest
15
16     steps:
17       # Checkout the merged PR code
18       - name: Checkout code
19         uses: actions/checkout@v3
20
21       # Log in to Docker Hub
22       - name: Log in to Docker Hub
23         uses: docker/login-action@v2
24         with:
25           username: ${{ secrets.DOCKER_USERNAME }}
26           password: ${{ secrets.DOCKER_PASSWORD }}
27
28       # Extract the GitHead Tag (Short SHA)
29       - name: Get commit SHA
30         id: get_sha
31         run: echo "sha_short=$(echo ${{ github.event.pull_request.head.sha }} | cut -c1-7)" >> $GITHUB_ENV
32
33       # Pull the current 'latest' image and tag it as 'stable'
34       - name: Retag latest as Stable
35         run: |
36           docker pull ${{ secrets.DOCKER_USERNAME }}/be-mern-server:latest
37           docker tag ${{ secrets.DOCKER_USERNAME }}/be-mern-server:latest ${{ secrets.DOCKER_USERNAME }}/be-mern-server:stable
38           docker push ${{ secrets.DOCKER_USERNAME }}/be-mern-server:stable

```

```

49 # Pull the new image, tag it as 'latest', and push it
50 - name: Tag and Push Latest
51   run: |
52     docker pull ${{ secrets.DOCKER_USERNAME }}/be-mern-server:${{ env.sha_short }}
53     docker tag ${{ secrets.DOCKER_USERNAME }}/be-mern-server:${{ env.sha_short }} ${{ secrets.DOCKER_USERNAME }}/be-mern-server:latest
54     docker push ${{ secrets.DOCKER_USERNAME }}/be-mern-server:latest
55
56 # Setup ssh key
57 - name: Set up SSH Key
58   run: |
59     echo "${{ secrets.EKS_CONTROL_PLANE_SSH_KEY }}" > eks_ssh_key
60     chmod 400 eks_ssh_key
61
62 # SSH into the server and restart the deployment
63 - name: Restart Deployment on EKS
64   run: |
65     ssh -o StrictHostKeyChecking=no -i eks_ssh_key ec2-user@${{ secrets.EKS_CONTROL_PLANE_IP_ADDRESS }} << 'EOF'
66     sudo su - << 'TIMER_EOF'
67     kubectl rollout restart deployment rmit-store-server-canary
68 EOF

```

- The stable and latest image is now available in Dockerhub

Containerized Microservices

Our applications are separated into 3 parts, a frontend, backend and a database. While MongoDB is already used as a separate service hosted on the mongo cloud itself, the frontend and backend are baked into Docker images for containerization.

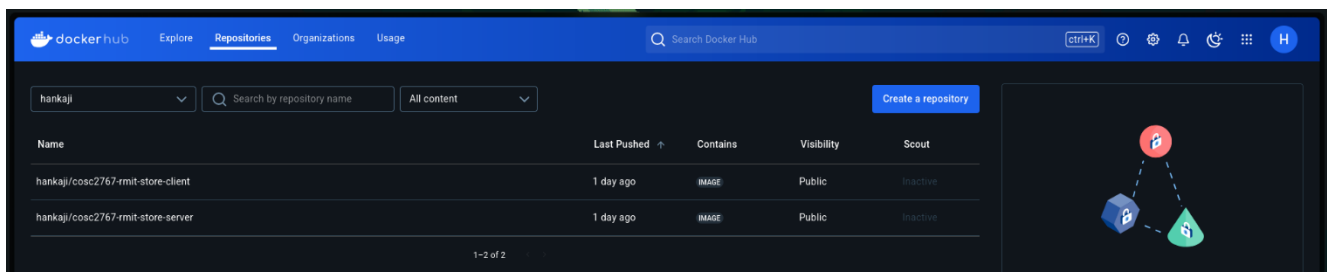


Figure 29. DockerHub Repo

These images are then pulled down by our Kubernetes deployment and deployed automatically into pods and clusters by Kubernetes. Furthermore, using EKS by AWS, we can separate our clusters into different worker group nodes, meaning that each deployment that contains a group of service clusters can run in different elastic containers, ensuring that in case of instability and any of our services goes down, it will not affect the other services and Kubernetes can easily replace it with another running services.

Node groups (3) [Info](#)

Edit

Delete

Add node group

Node groups implement basic compute scaling through EC2 Auto Scaling groups.

	Group name	Desired size	AMI release version	Launch template	Status
<input type="radio"/>	client-worker-group	2	1.31.3-20250103	-	<div><div></div>Active</div>
<input type="radio"/>	monitoring-group	2	1.31.3-20250103	-	<div><div></div>Active</div>
<input type="radio"/>	server-worker-group	2	1.31.3-20250103	-	<div><div></div>Active</div>

Figure 30. EKS Node Group

```

kubernetes/application/server-deployment.yml
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: rmit-store-server-pod
5   labels:
6     name: rmit-store-server-pod
7     app: rmit-store-server
8     version: stable
9   annotations:
10    prometheus.io/scrape: "true"
11    prometheus.io/path: "/metrics"
12    prometheus.io/port: "3000"
13 spec:
14   replicas: 2
15   selector:
16     matchLabels:
17       app: rmit-store-server
18       version: stable
19   template:
20     metadata:
21       labels:
22         app: rmit-store-server
23         name: rmit-store-server-pod
24         version: stable
25     annotations:
26       prometheus.io/scrape: "true"
27       prometheus.io/path: "/metrics"
28       prometheus.io/port: "3000"
29   spec:
30     containers:
31       - name: rmit-store-server-pod
32         image: localhost/be-mern-server:stable
33         imagePullPolicy: Always
34         ports:
35           - containerPort: 3000
36     nodeSelector:
37       eks.amazonaws.com/nodegroup: server-worker-group

```

By using Kubernetes' deployment service, we can dynamically scale our application services at runtime without needing to rewrite configurations each time, both horizontally and vertically. Scaling is straightforward and can be achieved with a few terminal commands. For horizontal scaling, we use the `kubectl scale` command to adjust the number of pods allocated to a deployment, enabling even distribution of traffic across multiple pods and reducing service workloads. For vertical scaling, the `kubectl resources` command allows us to allocate specific system resources (e.g., CPU, memory, and disk space) to deployments or individual clusters. When application demand decreases, we can scale down vertically to conserve resources and reduce costs, avoiding unnecessary resource allocation

Figure 31. Server Deployment Configuration

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/rmit-store-client-canary	0/0	0	0	14h
deployment.apps/rmit-store-client-pod	1/1	1	1	14h
deployment.apps/rmit-store-server-canary	0/0	0	0	14h
deployment.apps/rmit-store-server-pod	20/20	20	20	14h

Figure 32. Deployment list in Kubernetes

To enhance reliability, we run our Kubernetes clusters across multiple subnets, ensuring our services are available in multiple regions. This allows consumers from different parts of the world to experience reliable connectivity to our application without encountering network lag caused by regional restrictions.

Network environment info
 Configure cluster VPC settings

Subnets info
 Choose the subnets in your VPC where the control plane may place elastic network interfaces (ENIs) to facilitate communication with your cluster. To create a new subnet, go to the corresponding page in the [VPC console](#).

Select subnets

subnet-0a3b69b43860a1ab1
 us-east-1a 172.31.0.0/20

subnet-036298c633032c3df
 us-east-1b 172.31.80.0/20

subnet-07801c3051b8c3e37
 us-east-1c 172.31.16.0/20

Clear selected subnets

Additional security groups - optional info
 EKS automatically creates a cluster security group on cluster creation to facilitate communication between worker nodes and control plane. Optionally, choose additional security groups to apply to the EKS-managed Elastic Network Interfaces that are created in your control plane subnets. To create a new security group, go to the corresponding page in the [VPC console](#).

Select security groups

sg-0e47dce53a8e2989a | eks-security

Clear selected security groups

Figure 33. EKS Network Environment

Configuration management

Configurations are managed in two different ways. For our application, including the frontend and backend, we use environment variables. These variables are provided either through the Kubernetes container's environment settings or directly via a `dotenv` (.env) file. The configurations are applied at build time, so whenever changes are made to the application, a deployment restart is required for the new configurations to take effect.


```

server/.env
1 PORT=3000
2 # Local MongoDB URI Example
3 MONGO_URI_PROD=mongodb+srv://admin:123@cluster0.3vynp.mongodb.net/?retryWrites=true&w=majority&appName=Cluster0
4 MONGO_URI_TEST=mongodb://localhost:27017/rmit_ecommerce #Local MongoDB
5 appName=Cluster0
6 JWT_SECRET=my_secret_string
7 CLIENT_URL=http://rmit-store-client-server
8 BASE_API_URL=api
9 NODE_ENV=test

```

Figure 34. Dotenv configuration

```

kubernetes/application/client-deployment.yml
21 spec:
22   containers:
23     - name: rmit-store-client-pod
24       image: locoioioi/cosc2767-rmit-store-client:stable
25       imagePullPolicy: Always
26       ports:
27         - containerPort: 80
28       env:
29         - name: API_URL
30           value: /api
31   nodeSelector:
32     eks.amazonaws.com/nodegroup: client-worker-group

kubernetes/grafana.yml
43 containers:
44   - name: grafana
45     env:
46       - name: GF_SMTP_ENABLED
47         value: "true"
48       - name: GF_SMTP_HOST
49         value: "smtp.gmail.com:587"
50       - name: GF_SMTP_USER
51         value: "dtlamdevil@gmail.com"
52       - name: GF_SMTP_PASSWORD
53         value: "cuty rniy libv brdy"
54       - name: GF_SMTP_FROM_ADDRESS
55         value: "grafana@gmail.com"
56       - name: GF_SMTP_FROM_NAME
57         value: "Grafana"
58       - name: GF_SMTP_SKIP_VERIFY
59         value: "true"
60       image: grafana/grafana-enterprise
61       imagePullPolicy: IfNotPresent

```

Figure 35. Kubernetes deployment env

For other services like Prometheus, we host a ConfigMap and allow the service to use its configuration by connecting to the ConfigMap through a volume mount. Same as before, whenever changes are made to the ConfigMap, an application of Configmap and deployment restart is required for the new configurations to take effect.

```

kubernetes/prometheus-config.yml
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: prometheus-server-conf
5   namespace: monitoring
6 data:
7   prometheus.yml: |
8     global:
9       scrape_interval: 2m
10      evaluation_interval: 2m
11      body_size_limit: "0"
12      sample_limit: 0
13      label_limit: 0
14      label_name_length_limit: 0
15      label_value_length_limit: 0
16      target_limit: 0
17      scrape_configs:
18        - job_name: 'prometheus'
19          static_configs:
20            - targets: ['localhost:9090']
21        - job_name: 'node-exporter'
22          static_configs:
23            - targets: ['node-exporter.default.svc.cluster.local:9100']
24        - job_name: 'Store server'
25          kubernetes_sd_configs:
26            - role: pod
27              namespaces:
28                names:
29                  - default
30          relabel_configs:
31            - source_labels: [__meta_kubernetes_pod_label_version]
32              target_label: version
33            - source_labels: [__meta_kubernetes_pod_label_app]
34              target_label: app
35            - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_scrape]
36              action: keep
37              regex: true
38            - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_path]
39              action: replace
40              target_label: __metrics_path__
41            - source_labels: [__meta_kubernetes_pod_ip, __meta_kubernetes_pod_annotation_prometheus_io_port]
42              action: replace
43              target_label: __address__
44              regex: (.+)
45              separator: ":"

```

Figure 36. Configmap for Prometheus

```

kubernetes/prometheus-deployment.yml
15 spec:
16   serviceAccountName: prometheus
17   containers:
18   - name: prometheus
19     image: prom/prometheus
20     ports:
21     - containerPort: 9090
22     volumeMounts:
23     - name: config-volume
24       mountPath: /etc/prometheus
25   volumes:
26   - name: config-volume
27     configMap:
28       name: prometheus-server-conf
29       defaultMode: 420
30   nodeSelector:
31     eks.amazonaws.com/nodegroup: monitoring-group

```

Figure 37. Prometheus volume mount

Orchestration

As mentioned before, we use EKS, a Kubernetes managed by AWS, to deploy our applications using containerized services, which are pushed to Docker Hub.

Furthermore, by using Amazon EKS (Elastic Kubernetes Service), we can separate our clusters into different worker node groups. This allows each deployment, comprising a group of service clusters, to run in separate elastic containers. As a result, if instability occurs and one of our services goes down, it will not impact other services. Kubernetes can seamlessly replace the affected service with another running instance, ensuring minimal disruption.

Node groups (3) Info

Node groups implement basic compute scaling through EC2 Auto Scaling groups.

[Edit](#) [Delete](#) [Add node group](#)

Group name	Desired size	AMI release version	Launch template	Status
client-worker-group	2	1.31.3-20250103	-	Active
monitoring-group	2	1.31.3-20250103	-	Active
server-worker-group	2	1.31.3-20250103	-	Active

Figure 38. EKS Node groups

Nodes (6) Info

Filter Nodes by property or value

< 1 2 >

Node name	Instance type	Compute	Managed by	Created	Status
ip-172-31-13-73.ec2.internal	t2.medium	Node group	client-worker-group	Created 9 hours ago	Ready
ip-172-31-17-40.ec2.internal	t2.medium	Node group	monitoring-group	Created 9 hours ago	Ready
ip-172-31-19-135.ec2.internal	t2.medium	Node group	server-worker-group	Created 9 hours ago	Ready
ip-172-31-2-186.ec2.internal	t2.medium	Node group	monitoring-group	Created 9 hours ago	Ready
ip-172-31-30-217.ec2.internal	t2.medium	Node group	client-worker-group	Created 9 hours ago	Ready

Figure 39. EKS Nodes

```

kubernetes/application/server-deployment.yml
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: rmit-store-server-pod
5   labels:
6     name: rmit-store-server-pod
7     app: rmit-store-server
8     version: stable
9   annotations:
10    prometheus.io/scrape: "true"
11    prometheus.io/path: "/metrics"
12    prometheus.io/port: "3000"
13 spec:
14   replicas: 2
15   selector:
16     matchLabels:
17       app: rmit-store-server
18       version: stable
19   template:
20     metadata:
21       labels:
22         app: rmit-store-server
23       name: rmit-store-server-pod
24       version: stable
25     annotations:
26       prometheus.io/scrape: "true"
27       prometheus.io/path: "/metrics"
28       prometheus.io/port: "3000"
29   spec:
30     containers:
31       - name: rmit-store-server-pod
32         image: locolai/be-mern-server:stable
33         imagePullPolicy: Always
34         ports:
35           - containerPort: 3000
36     nodeSelector:
37       eks.amazonaws.com/nodegroup: server-worker-group

```

Figure 40. Server deployment configuration

The deployment of each service is managed through a YAML file, which instructs Kubernetes on the actions to take. To simplify this process, Kubernetes provides a Deployment resource type, which we use for the various services in our application. This resource type offers notable features, such as the ability to dynamically allocate system resources and adjust the number of clusters running for each service. This ensures flexibility in scaling our application to meet changing demands.

By specifying the number of replicas, we can tell Kubernetes how many pods to run for a desired service. This number can also be adjusted at runtime using the `kubectl scale` command. Similarly, resource allocations for each deployment can be modified dynamically as needed.

However, running pods alone are not nothing if they cannot be accessed. To address this, we use the Service resource in Kubernetes to expose our application to consumers. The type of Service we use depends on the service that we want to run. For example, for our front-end service, where end users interact directly, we use the Load Balancer service type. This allows Kubernetes to automatically manage internal IP addresses and ports, exposing the application to the outside world with an external IP address, which is far more user-friendly than a random sequence of numbers.

For other types, such as the backend or Node Exporter, it is unnecessary to expose them to the outside world since only internal applications and us developers need to interact with these services. In such cases, we use the Cluster IP service type, which restricts connection to internal communication within the cluster.

We connect with these services using the specific DNS domains supplied by Kubernetes, which are formed from the service name. For example, if our backend service is called `rmit-store-server-service`, we can access it internally via <http://rmit-store-server-service:<port>> where port is what we specified in deployment and service yaml files.

<pre> kubernetes/application/client-canary-deployment.yml 34 apiVersion: v1 35 kind: Service 36 metadata: 37 name: rmit-store-client-service 38 labels: 39 name: rmit-store-client-service 40 app: rmit-store-client 41 spec: 42 ports: 43 - port: 80 44 targetPort: 80 45 type: LoadBalancer 46 selector: 47 app: rmit-store-client </pre>	<pre> kubernetes/application/server-deployment.yml 39 apiVersion: v1 40 kind: Service 41 metadata: 42 name: rmit-store-server-service 43 labels: 44 name: rmit-store-server-service 45 app: rmit-store 46 spec: 47 ports: 48 - port: 3000 49 targetPort: 3000 50 selector: 51 app: rmit-store-server </pre>
---	---

Figure 41. LoadBalancer & ClusterIP Service

IV. Pipeline Advanced Component

8. Infrastructure as Code (IaC)

This is a service that helps us model and set up AWS resources and provides more time for applications operating in AWS rather than dealing with resource management [4]. It is used to bring up the test environment that simulates the production environment then we can shut down and eliminate that environment after testing for cost efficiency.

```

AWSTemplateFormatVersion: "2010-09-09"
Resources:
  TestEnvEC2Instance:
    Type: "AWS::EC2::Instance"
    Properties:
      InstanceType: "t2.micro"
      KeyName: "vockey" # The key used in the Jenkinsfile credential "vockey"
      ImageId: "ami-05576a079321f21f8" # Amazon Linux 2 AMI
      SecurityGroupIds:
        - !Ref TestEnvSecurityGroup
      UserData:
        Fn::Base64: !Sub |
          #!/bin/bash
          yum update -y
          yum install -y docker
          service docker start
          usermod -s /bin/bash ec2-user "usermod": Unknown word,
          amazon-linux-extras enable nginx1
          yum install -y nodejs npm
          curl -fsSL https://deb.nodesource.com/setup_18.x | bash -
          yum install -y mongodb-org
          systemctl start mongod
          systemctl enable mongod

  TestEnvSecurityGroup:
    Type: "AWS::EC2::SecurityGroup"
    Properties:
      GroupDescription: "Enable SSH and HTTP access"
      SecurityGroupIngress:
        - IpProtocol: "tcp"
          FromPort: "22"
          ToPort: "22"
          CidrIp: "0.0.0.0/0"
        - IpProtocol: "tcp"
          FromPort: "80"
          ToPort: "80"
          CidrIp: "0.0.0.0/0"
        - IpProtocol: "tcp"
          FromPort: "3000"
          ToPort: "3000"
          CidrIp: "0.0.0.0/0"

Outputs:
  PublicIP:
    Description: Public IP address of the EC2 instance
    Value: !GetAtt TestEnvEC2Instance.PublicIp

```

Figure 42 AWS CloudFormation template

This template is mainly focused on set up an environment for testing purpose and it will be configured as below:

- **Type:** AWS::EC2::Instance
- **Properties:**
 - **Instance type:** it will specify the type of instance that will be used for the EC2 instance, for example, figure 21 will create t2.micro instance.
 - **KeyName:** This will specify what private key will be used for connecting to the instance.
 - **SecurityGroupIds:** Associates the instance with TestEnvSecurityGroup for security configurations.
 - **UserData:** This will specify the script will be executed after the instance is up and running.

9. Monitoring and Alerting

a. Prometheus

Prometheus is an open-source application monitoring system that collects metrics exposed by various parts of our application. To set up Prometheus, we first configure our application to expose metrics that Prometheus can scrape to monitor application health and performance. Next, we configure Prometheus to scrape these metrics by specifying the endpoints and the scrape interval in a YAML configuration file, as shown below:

```

kubernetes/prometheus-config.yml
8 global:
9   scrape_interval: 2m
10  evaluation_interval: 2m
11  body_size_limit: "0"
12  sample_limit: 0
13  label_limit: 0
14  label_name_length_limit: 0
15  label_value_length_limit: 0
16  target_limit: 0
17 scrape_configs:
18   - job_name: 'prometheus'
19     static_configs:
20       - targets: ['localhost:9090']
21   - job_name: 'node-exporter'
22     static_configs:
23       - targets: ['node-exporter.default.svc.cluster.local:9100']
24   - job_name: 'Store_server'
25     kubernetes_sd_configs:
26       - role: pod
27         namespaces:
28           names:
29             - default
30     relabel_configs:
31       - source_labels: [__meta_kubernetes_pod_label_version]
32         target_label: version
33       - source_labels: [__meta_kubernetes_pod_label_app]
34         target_label: app
35       - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_scrape]
36         action: keep
37         regex: true
38       - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_path]
39         action: replace
40         target_label: __metrics_path__
41       - source_labels: [__meta_kubernetes_pod_ip, __meta_kubernetes_pod_annotation_prometheus_io_port]
42         action: replace
43         target_label: __address__
44         regex: (.+)
45         separator: ":"

```

Figure 43. Prometheus config

The global scrape and evaluate interval instruct Prometheus that it needs to scrape and evaluate every 2 minutes intervals, applying to all jobs.

We have two main jobs to monitor: 'Node Exporter' and 'Store Server.' For the Node Exporter, we deploy it using the Daemon Set resource provided by Kubernetes, which ensures it runs on all nodes within a namespace (default in this case). These Node Exporters monitor system resources and usage on each node, exposing the metrics data at `http://<node-ip>:9100/metrics`.

<pre> kubernetes/rbac-prometheus.yml 2 apiVersion: rbac.authorization.k8s.io/v1 3 kind: ClusterRole 4 metadata: 5 name: prometheus 6 rules: 7 - apiGroups: [""] 8 resources: 9 - pods 10 - services 11 - endpoints 12 - nodes 13 - namespaces 14 verbs: ["get", "list", "watch"] 15 --- 16 apiVersion: rbac.authorization.k8s.io/v1 17 kind: ClusterRoleBinding 18 metadata: 19 name: prometheus-binding 20 namespace: monitoring 21 roleRef: 22 apiGroup: rbac.authorization.k8s.io 23 kind: ClusterRole 24 name: prometheus 25 subjects: 26 - kind: ServiceAccount 27 name: prometheus 28 namespace: monitoring 29 --- 30 apiVersion: v1 31 kind: ServiceAccount 32 metadata: 33 name: prometheus 34 namespace: monitoring </pre>	<pre> kubernetes/node-exporter-daemonset.yml 1 apiVersion: apps/v1 2 kind: DaemonSet 3 metadata: 4 labels: 5 app.kubernetes.io/component: exporter 6 app.kubernetes.io/name: node-exporter 7 name: node-exporter 8 spec: 9 selector: 10 matchLabels: 11 app.kubernetes.io/component: exporter 12 app.kubernetes.io/name: node-exporter 13 template: 14 metadata: 15 labels: 16 app.kubernetes.io/component: exporter 17 app.kubernetes.io/name: node-exporter 18 spec: 19 containers: 20 - args: 21 - --path.sysfs=/host/sys 22 - --path.rootfs=/host/root 23 - --no-collector.wifi 24 - --no-collector.human 25 - --collector.filesystem.ignored-mount-points='[/dev proc sys var/lib/docker/.+ var/lib/kubelet/pods/.+](/)?' 26 - --collector.netclass.ignored-devices='(veth.+)' 27 name: node-exporter 28 image: prom/node-exporter 29 ports: 30 - containerPort: 9100 31 protocol: TCP 32 resources: 33 limits: 34 cpu: 250m 35 memory: 180Mi 36 requests: 37 cpu: 102m 38 memory: 180Mi 39 volumeMounts: 40 - mountPath: /host/sys 41 mountPropagation: HostToContainer 42 name: sys 43 readOnly: true 44 - mountPath: /host/root 45 mountPropagation: HostToContainer 46 name: root 47 readOnly: true 48 volumes: 49 - hostPath: 30 path: /sys 31 name: sys 32 - hostPath: 33 path: / 34 name: root </pre>
---	--

Figure 44. Prometheus & Node exporter configuration

However, since our Prometheus deployment and Node Exporter are in different namespaces, Prometheus does not have sufficient permissions to access pods and resources in other namespaces. To address this, we need to create RBAC (Role-Based Access Control) permissions by defining Cluster Roles and Cluster Role Bindings. Additionally, we configure a custom Service Account for Prometheus, granting it the necessary permissions to access the required pods and namespaces.

```
server/index.js
38 // Middleware to enable prometheus scraping
39 const metricsMiddleware = promBundle({
40   includeMethod: true,
41   includePath: true,
42   includeStatusCode: true,
43   includeUp: true,
44   customLabels: {
45     project_name: "Rmit-store",
46     project_type: "Rmit-server-metrics",
47   },
48   promClient: {
49     collectDefaultMetrics: {},
50   },
51 });
52
53 app.use(metricsMiddleware);
```

Figure 45. Prom Client configuration

For the Store Server, we use prom-client and express-prom-bundle to expose metrics about system usage and HTTP request rates. These metrics are made available at `http://<node-ip>:3000/metrics`. However, instead of using a static endpoint to scrape data, we leverage Kubernetes' service discovery feature to dynamically scrape all service pods. This ensures that Prometheus can correctly gather data even when new clusters are added. Additionally, this approach automatically labels the data based on pod labels, allowing us to easily query and differentiate metrics for canary and stable deployments.

b. Grafana

Grafana is a powerful tool for visualizing data queried from a data source, such as Prometheus in our case. It allows us to import pre-built dashboards or create custom visualizations to monitor data trends over time through graphs and other visual formats. This is particularly useful for identifying application instabilities or unexpected traffic spikes, enabling us to scale the application appropriately without over-provisioning resources.

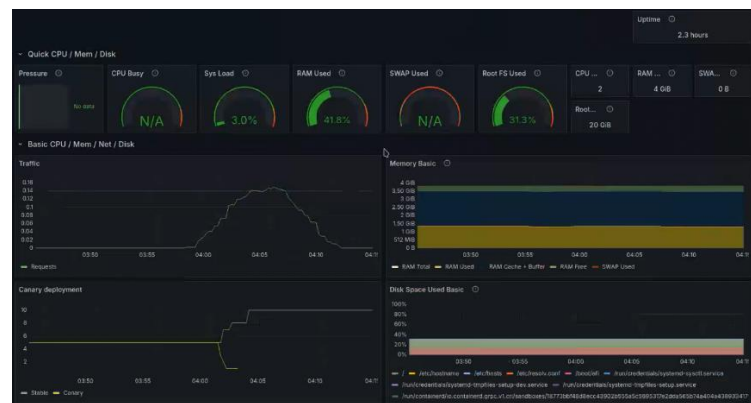


Figure 46. Grafana Dashboard

First, we demonstrate a traffic spike in our application by querying the HTTP request count from Prometheus, which shows the number of requests being made to our application. As illustrated in the figure above, the request count spikes significantly before gradually decreasing over the next few minutes. This behavior was simulated using curl to send requests to the application every 0.5 seconds for a total of 100 requests.

There is also a visualization for the canary deployment, which displays the number of services currently running. As shown in the figure above, the server service of our application starts with both the canary and stable versions running. Over time, the canary version gradually scales down to 0 instances, while the stable version scales up to 10 running instances.

Moreover, we have configured alert rules to notify us when certain thresholds are exceeded. For example, if the traffic spike surpasses a threshold of 0.1, an email alert is sent, notifying us of the high request rate in our application. This allows us to take quick action before the application becomes stressed. Similarly, an alert is triggered if RAM usage exceeds 1.5GB, informing us of high memory consumption so we can address the issue promptly.

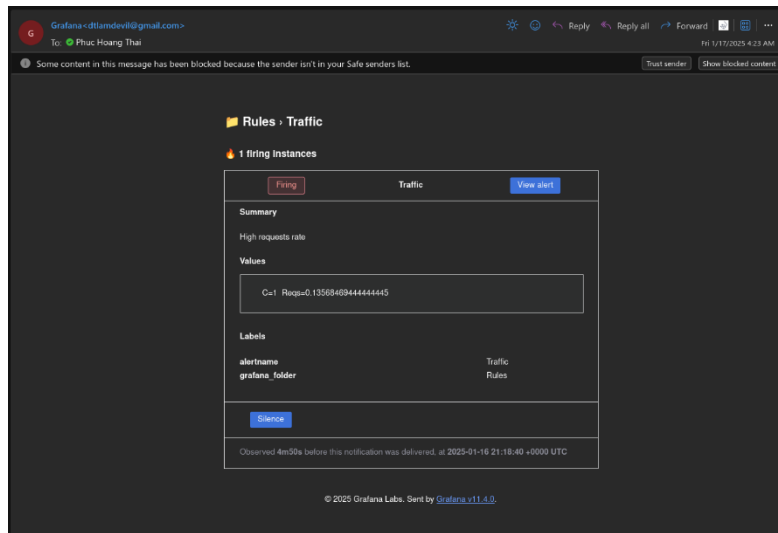


Figure 47. Grafana Email Notification

10.Canary Deployment

In our CI/CD application, we will utilize the canary strategy which will roll out the latest version of the application gradually to only a small subset of users and keep the majority of users stay on the stable deployed version. Once there are no issues with the canary releases, it will be available to broader audiences and a smaller proportion of already-deployed versions.

In the pipeline, for client and server instance, each instance will have the deployment for stable and canary version. As the configuration for the canary deployment strategies is similar among the front-end and back-end instance, only the code for the client are explained.

Stable Version Deployment

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: rmit-store-client-pod
5    labels:
6      name: rmit-store-client-pod
7      app: rmit-store-client
8      version: stable
9  spec:
10   replicas: 2
11   selector:
12     matchLabels:
13       app: rmit-store-client
14       version: stable
15   template:
16     metadata:
17       labels:
18         name: rmit-store-client-pod
19         app: rmit-store-client
20         version: stable
21     spec:
22       containers:
23         - name: rmit-store-client-pod
24           image: locoioloi/cosc2767-rmit-store-client:stable
25           imagePullPolicy: Always
26           ports:
27             - containerPort: 80
28           env:
29             - name: API_URL
30               value: /api
31       nodeSelector:
32         eks.amazonaws.com/nodegroup: client-worker-group
33 ---
34 apiVersion: v1
35 kind: Service
36 metadata:
37   name: rmit-store-client-service
38   labels:
39     name: rmit-store-client-service
40     app: rmit-store-client
41 spec:
42   ports:

```

Figure 48. Stable version deployment configuration file

In the stable version deployment configuration, This YAML configuration defines the deployment and service for the stable version of the rmit-store-client application, supporting canary deployment. The deployment, named 'rmit-store-client-pod', creates 2 replicas using the stable container image 'locoioloi/cosc2767-rmit-store-client:stable' from DockerHub, with the imagePullPolicy set to 'Always', ensuring the latest image is pulled during a 'kubectl rollout restart'. Pods are assigned to the 'client-worker-group' node group for efficient workload distribution when scaling vertically or horizontally and expose port '80', with an environment variable 'API_URL' set to '/api'.

The associated service, 'rmit-store-client-service', uses a 'LoadBalancer' to route traffic from the LoadBalancer IP to the pods on port '80', enabling balanced external access. This setup facilitates smooth updates, efficient traffic routing, and scalability during deployments.

Canary Version Deployment

```
cosc2767-assignment-2-group-2024c-pall - client-canary-deployment.yml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: rmit-store-client-canary
5    labels:
6      name: rmit-store-client-canary
7      app: rmit-store-client
8      version: stable
9  spec:
10   replicas: 1
11   selector:
12     matchLabels:
13       app: rmit-store-client
14       version: stable
15   template:
16     metadata:
17       labels:
18         name: rmit-store-client-canary
19         app: rmit-store-client
20         version: stable
21     spec:
22       containers:
23         - name: rmit-store-client-canary
24           image: locoioloi/cosc2767-rmit-store-client:latest
25           imagePullPolicy: Always
26           ports:
27             - containerPort: 80
28           env:
29             - name: API_URL
30               value: /api
31       nodeSelector:
32         eks.amazonaws.com/nodegroup: client-worker-group
33 ---
34 apiVersion: v1
35 kind: Service
36 metadata:
37   name: rmit-store-client-service
38   labels:
39     name: rmit-store-client-service
40     app: rmit-store-client
41 spec:
42   ports:
43     - port: 80
44       targetPort: 80
45   type: LoadBalancer
46   selector:
47     app: rmit-store-client
48
49
```

Figure 49. Canary version deployment configuration file

For the latest deployment, the canary YAML configuration defines the deployment and service for the canary version of the rmit-store-client application, supporting canary deployment. The deployment, named 'rmit-store-client-canary', creates one replica using the latest container image 'locoioloi/cosc2767-rmit-store-client:latest' from DockerHub, with the imagePullPolicy set to 'Always', ensuring the latest image is pulled during a 'kubectl rollout restart'. The pod is assigned to the 'client-worker-group' node group to optimize workload distribution and exposes port '80', with an environment variable 'API_URL' set to '/api'. The associated service, 'rmit-store-client-service', uses a 'LoadBalancer' to route traffic from the LoadBalancer IP to the pods on port '80', enabling balanced external access. This setup facilitates testing the latest version of the application while directing minimal traffic to the canary deployment, ensuring reliability during deployments.

Stable & Canary Version Traffic Distribution

As both of the stable and canary deployment: 'rmit-store-client-pod', 'rmit-store-client-canary' are both attached to the 'rmit-store-client-service' under the Load Balancer. When access the Load Balancer IP, the weight of distribution of 2 versions are calculate as below:

- **n1 = 2**: Number of stable pod replicas for 'rmit-store-client-pod'
- **n2 = 1**: Number of canary pod replicas for 'rmit-store-client-canary'

$$\text{Traffic Distribution to Stable Version Deployment} = \frac{n1}{n1 + n2} \sim 66\%$$

$$\text{Traffic Distribution to Canary Version Deployment} \sim 33\%$$

Based on this mechanism of Kubernetes traffic splitting, to make the canary version accessible to broader user audiences, the command 'kubectl scale deployment [deployment_name] --replicas=[number]' are used to either scale up the number of canary pod or scale down the stable pod replicas.

V. Application Flow (Diagram)

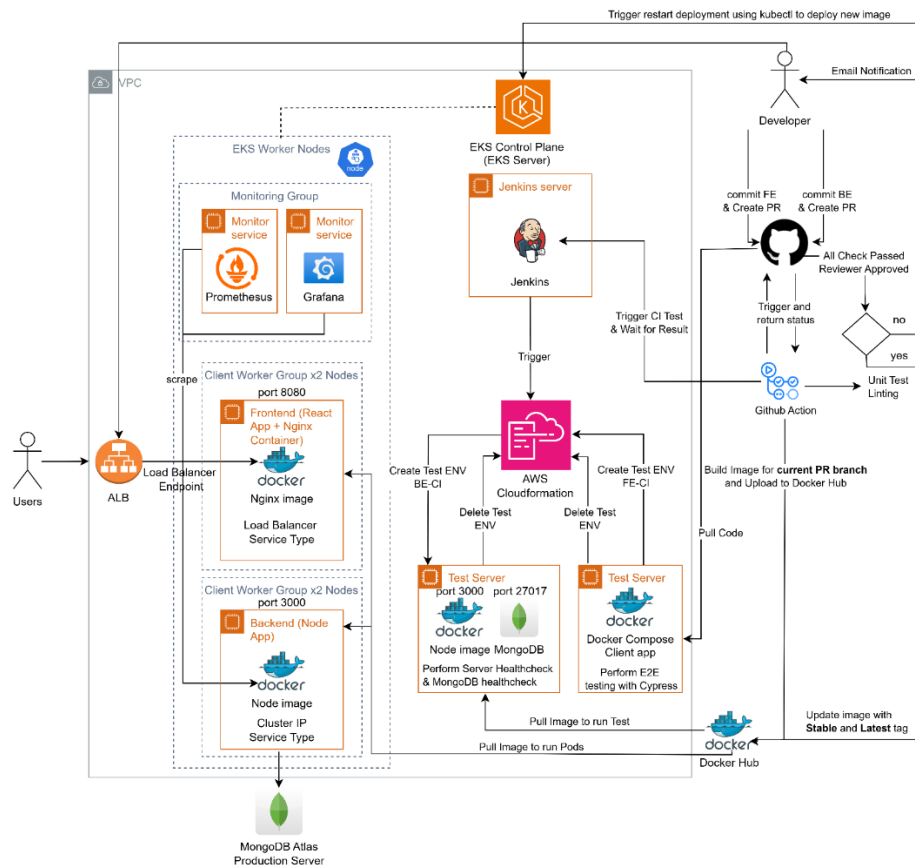


Figure 50. Application Flow Diagram

- GitHub Pull Request Merge to Main:** Code commit and push directly to Main branch is blocked, so the Developer will need to commit the code change to a feature branch, and create a Pull Request (PR) to merge that branch into Main branch. The PR will trigger a series of tests to ensure correctness of the code and require approval from other contributors to be able to merge.
- GitHub Action:** Depending on the code changes, a series of GitHub Action workflow will be triggered.
Unit Test and Integration Test will always be triggered for any changes to both server or client code, and required to be pass, if test cases is not provided, this test will skip with success.
Linting Check will also be triggered for any changes, lint check is not required to be pass in order to be able to merge, and serve as a recommendation for better coding styles and convention.
 - Back-end changes (/server/**): This will trigger jenkins-ci-pr.yaml workflow, which will create an image docker and a test environment for the current PR branch.
 - Front-end changes (/client/**): This will trigger jenkins-ci-fe.yaml workflow, which will create a test environment and conduct end-to-end testing for client.

After a PR is merged into Main, there will also be a workflow to build and update docker image to the newest tag in Docker Hub, then trigger Kubernetes deployment on the EKS server.

- Jenkins:** GitHub Action will trigger the corresponding jobs be-ci or fe-ci to create the test environment, install the application and perform HealthCheck/testing.
- AWS CloudFormation:** Use the CloudFormation template stored in the GitHub repository to create a new EC2 instance every time a new test environment is created. The application will be installed and tested in this environment, then the instance will be terminated after the test.
- Docker Hub:** Docker Hub is used to store Docker Images built from the pipeline for testing and deployment to EKS.

6. **AWS Elastic Kubernetes Service (EKS):** Kubernetes cluster used to deploy the application pulled from Docker Hub, and the monitoring service. This cluster have 3 groups:
 - a. Client Worker Group
 - b. Server Worker Group
 - c. Monitoring Group
7. **MongoDB Atlas:** The cloud-hosted MongoDB is used for production environment for the application to ensure its isolation with the test environment, and reduced management effort needed to connect to the backend running in EKS.
8. **Prometheus & Grafana:** Prometheus will scrape data and hardware usage information from the Kubernetes cluster and the application. Grafana will use that as the data source and visualize them into a dashboard.
9. **Application Load Balancer (ALB):** This is the endpoint for the user to access and interact with the application. This load balancer will ensure fair distribution and spread the load across nodes in the Kubernetes Cluster, it also serves the purpose of canary deployment type.
10. These worker nodes are managed under a Network Load Balancer that distributes incoming network traffic across multiple servers. This ensures no single server bears too much demand.

VI. Known Bugs/ Problems

- **Manual Canary Deployment Visibility:** The current process for canary deployment requires manually scaling traffic to new version pods to evaluate the updated deployment. This manual intervention increases operational overhead and slows down the deployment process. It also makes it more challenging to compare the performance and reliability of the new version against the stable one in real-time.
- **Vertical Scaling Monitor:** There is a lack of real-time metrics to monitor the computational resources utilized by containers. Without these metrics, the system is unable to identify and respond to resource bottlenecks or overutilization effectively. This limitation prevents proactive management of container performance and could lead to service degradation during high resource usage or unexpected traffic spikes.

VII. Pipeline Self-Evaluation

Completing this project provided our team with hands-on experience in building a secure, scalable, and automated DevOps pipeline while adhering to industry-standard security practices. Guided by insights from industry guest speakers, we implemented a private GitHub repository to safeguard source code, enforced approval-based pull requests to protect the main branch, and used environment variables to securely manage sensitive information. These measures ensured robust access control, reduced vulnerabilities, and adhered to best practices.

Our team believes this solution is the most appropriate because it combines automation, scalability, and reliability to address integration and deployment challenges. Automated CI/CD workflows with Jenkins ensure early error detection and consistent delivery, while Docker and Kubernetes provide portability, reliability, and efficient orchestration. AWS services like CloudFormation and Elastic Load Balancing enable scalable, consistent infrastructure management.

Automation

The pipeline is designed to automate key stages of software delivery, from code integration to testing and deployment. GitHub Actions, in conjunction with Jenkins, triggers workflows immediately upon code changes or pull requests. AWS CloudFormation provisions infrastructure automatically, while Kubernetes orchestrates container deployments, making the release process less dependent on human intervention. This extensive automation minimizes manual errors, accelerates feedback loops, and ensures consistency throughout the software lifecycle.

Efficiency

By automating repetitive tasks, the pipeline optimizes resource usage and developer time. Changes can be tested and validated shortly after they are committed, enabling faster feedback and reducing the time spent on identifying and resolving issues. Containerization with Docker further speeds up deployments by standardizing environments, and infrastructure as code via AWS CloudFormation prevents the overhead associated with manual environment setup or teardown.

Reliability

Reliability is significantly increased through automated testing and monitoring at multiple levels. Jenkins ensures each build passes standardized tests, while Kubernetes and Nginx support zero-downtime deployments and provide resilience through rolling updates. Prometheus and Grafana supply real-time system metrics and alerts, allowing teams to identify and address anomalies before they impact end-users. Rollback capabilities in Kubernetes enable a swift return to a stable version if a deployment encounters critical issues, further enhancing reliability.

Scalability

Scalability is a key focus of this setup. Elastic Load Balancing distributes incoming traffic across multiple servers or containers, mitigating performance bottlenecks. Kubernetes automatically scales containers based on metrics such as CPU and memory usage, while MongoDB Atlas accommodates growing data needs without manual database administration. These features collectively ensure that the RMIT E-commerce platform can handle increased traffic and user demands with minimal operational overhead.

Repeatability

The use of containerization and infrastructure as code promotes a highly repeatable deployment process. Docker images guarantee that each environment runs an identical stack, reducing the likelihood of environment-specific defects. With AWS CloudFormation templates, it is straightforward to replicate production-like environments for development and testing. This uniformity eliminates inconsistencies between environments and makes the entire pipeline—from code commit to deployment—easily reproducible for each new release cycle.

VIII. Conclusion

Completing this project has given our team valuable hands-on experience in designing and implementing a comprehensive DevOps pipeline using modern technologies. The project provided an in-depth understanding of the importance of automation, scalability, and reliability in delivering robust software solutions. Our pipeline automates critical processes such as code integration, build, testing, and deployment, significantly reducing manual intervention and potential errors. Through continuous integration (CI), every new pull request triggers an automated build and test process in Jenkins, ensuring early detection of errors and maintaining synchronization across the repository. Continuous delivery (CD) further enables seamless deployment to testing and production environments, improving release management flexibility.

Leveraging AWS CloudFormation for automated test environment management, the pipeline ensures consistency between testing and production environments while maintaining cost efficiency by dynamically tearing down resources when not in use. Docker simplifies the deployment process by packaging applications into lightweight containers, ensuring consistent performance across development, testing, and production. Kubernetes orchestrates these containerized applications, enabling features like self-healing, canary deployments, and rapid rollbacks, enhancing agility and minimizing downtime. Traffic scalability is achieved with Elastic Load Balancing, which dynamically adjusts to maintain optimal performance under varying traffic loads.

Additionally, the integration of Prometheus and Grafana ensures real-time monitoring and alerting, allowing the team to proactively address anomalies and performance issues before they impact users. The hands-on use of these technologies—Jenkins, Docker, AWS, Kubernetes, Prometheus, and Grafana—has provided practical DevOps knowledge and reinforced the importance of infrastructure-as-code, containerization, and monitoring in modern software development. Beyond technical skills, the project also honed our critical thinking, problem-solving, and adaptability, as we navigated challenges and developed efficient automation solutions. This project has greatly enriched our understanding of building scalable, reliable, and automated systems, preparing us for future challenges in the field of DevOps.

IX. Project Responsibility

Name	Role	Task	Contribution
Do Tung Lam	EKS Engineer, Data Engineer	Set up EKS, define nodes, pods, service, set up canary deployment strategies, set up Prometheus, Grafana for visualization	20%

Hoang Thai Phuc	Docker Engineer, Data Engineer	Set up Dockerfile, Docker compose for the application, set up Prometheus, Grafana for visualization	20%
Nguyen Quoc An	Tester, Jenkins Engineer	Creating and executing test cases, ensuring the availability of the website, Set up Jenkins Jobs	20%
Truong Quang Bao Loc	Tester, GitHub actions Engineer	Creating and executing test cases, ensuring the availability of the website, Set up GitHub actions Job	20%

X. Reference

- [1] GitHub, "GitHub," *Wikipedia*, Dec. 2023. [Online]. Available: <https://en.wikipedia.org/wiki/GitHub>. [Accessed: 17-Jan-2025].
- [2] GitHub, "GitHub Actions Documentation," *GitHub Docs*. [Online]. Available: <https://docs.github.com/en/actions>. [Accessed: 17-Jan-2025].
- [3] Jenkins, "Jenkins User Documentation," *Jenkins.io*. [Online]. Available: <https://www.jenkins.io/doc/>. [Accessed: 17-Jan-2025].
- [4] Amazon Web Services, "AWS CloudFormation User Guide," *Amazon AWS*, 2025. [Online]. Available: <https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/Welcome.html>. [Accessed: 17-Jan-2025].
- [5] Amazon Web Services, "Docker on AWS," *Amazon AWS*, 2025. [Online]. Available: https://aws.amazon.com/docker/?nc1=h_ls. [Accessed: 17-Jan-2025].
- [6] MongoDB, "MongoDB Atlas Tutorial," *MongoDB Resources*. [Online]. Available: <https://www.mongodb.com/resources/products/platform/mongodb-atlas-tutorial>. [Accessed: 17-Jan-2025].
- [7] NGINX, "Welcome to NGINX," *NGINX.org*. [Online]. Available: <https://nginx.org/en/>. [Accessed: 17-Jan-2025].
- [8] Amazon Web Services, "Elastic Load Balancing," *Amazon AWS*, 2025. [Online]. Available: <https://aws.amazon.com/elasticloadbalancing/>. [Accessed: 17-Jan-2025].
- [9] Kubernetes, "Kubernetes Documentation," *Kubernetes.io*. [Online]. Available: <https://kubernetes.io/>. [Accessed: 17-Jan-2025].
- [10] Prometheus, "Prometheus Monitoring System," *Prometheus.io*. [Online]. Available: <https://prometheus.io/>. [Accessed: 17-Jan-2025].
- [11] Grafana Labs, "Grafana – The open observability platform," *Grafana.com*. [Online]. Available: <https://grafana.com/>. [Accessed: 17-Jan-2025].

XI. Appendix